

Aim:

To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Steps:**1. Create a key pair.**

A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.

Name

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type Info
☒ RSA ☐ ED25519

Private key file format
☒ .pem
For use with OpenSSH
☐ .ppk
For use with PuTTY

Tags - *optional*
No tags associated with the resource.

You can add up to 50 more tags.

Key pairs (2) Info

<input type="checkbox"/>	Name	Type	Created	Fingerprint	ID
<input type="checkbox"/>	rook	rsa	2024/09/14 22:42 GMT+5:30	40:38:ad:9e:d0:9d:51:f4:f1:20:b0:...	key-04
<input type="checkbox"/>	ec2	rsa	2024/08/05 13:07 GMT+5:30	15:77:76:82:87:d2:40:e4:db:c7:2a:...	key-0b

The .pem file will be downloaded on your machine and will be required in the further steps.

2. Now we will create an EC2 Ubuntu instance. Select the key pair which you just created while creating this instance.

Instances (1) Info Last updated less than a minute ago

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availi
<input type="checkbox"/>	instance	i-051e99d82072f03cd	Running	t2.micro	2/2 checks passed	View alarms	ap-so

3. Now edit the inbound rules to allow ssh.

- ## 6. Docker installation:

We will be installing docker by using “sudo yum install docker -y”

```

Last metadata expiration check: 0:05:38 ago on Sat Sep 14 17:38:25 2024.
Dependencies resolved.
=====
Package                               Architecture      Version           Size
-----
Installing:
docker                                x86_64            25.0.6-1.amzn2023.0.2      44 M
Installing dependencies:
containerd                            x86_64            1.7.20-1.amzn2023.0.1      35 M
iptables-libs                         x86_64            1.8.8-3.amzn2023.0.2      401 k
iptables-nft                         x86_64            1.8.8-3.amzn2023.0.2      183 k
libcgroup                             x86_64            3.0-1.amzn2023.0.1         75 k
libnetfilter_conntrack               x86_64            1.0.8-2.amzn2023.0.2      58 k
libnftnl                             x86_64            1.0.1-19.amzn2023.0.2      30 k
libnftnl                             x86_64            1.2.2-2.amzn2023.0.2      84 k
pigz                                  x86_64            2.5-1.amzn2023.0.3         83 k
runc                                  x86_64            1.1.13-1.amzn2023.0.1      3.2 M
-----
Transaction Summary
-----
Install 10 Packages

Total download size: 84 M
Installed size: 317 M
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_64.rpm      4.1 MB/s | 401 kB  00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_64.rpm      6.8 MB/s | 183 kB  00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm           1.4 MB/s | 75 kB   00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64.rpm 3.1 MB/s | 58 kB   00:00
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_64.rpm         1.2 MB/s | 30 kB   00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rpm          2.0 MB/s | 84 kB   00:00
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm                1.4 MB/s | 83 kB   00:00
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm             15 MB/s | 3.2 MB   00:00
(9/10): containerd-1.7.20-1.amzn2023.0.1.x86_64.rpm       34 MB/s | 35 MB    00:01
(10/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rpm         32 MB/s | 44 MB    00:01
-----
Total
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction:
Preparing :
Installing : runc-1.1.13-1.amzn2023.0.1.x86_64              1/1
Installing : containerd-1.7.20-1.amzn2023.0.1.x86_64        2/10
Running scriptlet: containerd-1.7.20-1.amzn2023.0.1.x86_64 2/10
Installing : pigz-2.5-1.amzn2023.0.3.x86_64                 3/10
Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64           4/10
Installing : libnftnl-1.0.1-19.amzn2023.0.2.x86_64          5/10
Installing : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64 6/10
Downloading Packages:
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_64.rpm      4.1 MB/s | 401 kB  00:00
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_64.rpm      6.8 MB/s | 183 kB  00:00
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm           1.4 MB/s | 75 kB   00:00
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64.rpm 3.1 MB/s | 58 kB   00:00
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_64.rpm         1.2 MB/s | 30 kB   00:00
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rpm          2.0 MB/s | 84 kB   00:00
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm                1.4 MB/s | 83 kB   00:00
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm             15 MB/s | 3.2 MB   00:00
(9/10): containerd-1.7.20-1.amzn2023.0.1.x86_64.rpm       34 MB/s | 35 MB    00:01
(10/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rpm         32 MB/s | 44 MB    00:01
-----
Total
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction:
Preparing :
Installing : runc-1.1.13-1.amzn2023.0.1.x86_64              1/1
Installing : containerd-1.7.20-1.amzn2023.0.1.x86_64        2/10
Running scriptlet: containerd-1.7.20-1.amzn2023.0.1.x86_64 2/10
Installing : pigz-2.5-1.amzn2023.0.3.x86_64                 3/10
Installing : libnftnl-1.2.2-2.amzn2023.0.2.x86_64           4/10
Installing : libnftnl-1.0.1-19.amzn2023.0.2.x86_64          5/10
Installing : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64 6/10
Installing : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64       7/10
Installing : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64       8/10
Running scriptlet: iptables-nft-1.8.8-3.amzn2023.0.2.x86_64 8/10
Installing : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64       9/10
Running scriptlet: libcgroup-3.0-1.amzn2023.0.1.x86_64       10/10
Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64      10/10
Installing : docker-25.0.6-1.amzn2023.0.2.x86_64            10/10
Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
Verifying : containerd-1.7.20-1.amzn2023.0.1.x86_64          1/10
Verifying : docker-25.0.6-1.amzn2023.0.2.x86_64             2/10
Verifying : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64       3/10
Verifying : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64       4/10
Verifying : libcgroup-3.0-1.amzn2023.0.1.x86_64            5/10
Verifying : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64 6/10
Verifying : libnftnl-1.0.1-19.amzn2023.0.2.x86_64          7/10
Verifying : libnftnl-1.2.2-2.amzn2023.0.2.x86_64           8/10
Verifying : pigz-2.5-1.amzn2023.0.3.x86_64                  9/10
Verifying : runc-1.1.13-1.amzn2023.0.1.x86_64              10/10

Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64      docker-25.0.6-1.amzn2023.0.2.x86_64      iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
iptables-nft-1.8.8-3.amzn2023.0.2.x86_64    libcgroup-3.0-1.amzn2023.0.1.x86_64      libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
libnftnl-1.0.1-19.amzn2023.0.2.x86_64      libnftnl-1.2.2-2.amzn2023.0.2.x86_64      pigz-2.5-1.amzn2023.0.3.x86_64
runc-1.1.13-1.amzn2023.0.1.x86_64

Complete!
[ec2-user@ip-172-31-3-16 ~]$

```

7. Then to configure cgroup in a daemon json file we will run `cd /etc/docker` cat `<<EOF | sudo tee /etc/docker/daemon.json`
- ```
{
 "exec-opts": ["native.cgroupdriver=systemd"]
}
```

```
}
EOF
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker

[ec2-user@ip-172-31-3-16 ~]$ cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
 "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
{
 "exec-opts": ["native.cgroupdriver=systemd"]
}
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-3-16 docker]$ |
```

## 8. Kubernetes installation:

Search kubeadm installation on your browser and scroll down and select red hat-based distributions.

1. Set SELinux to `permissive` mode:

These instructions are for Kubernetes 1.31.

```
Linux in permissive mode (effectively disabling it)
setenforce 0
sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
This overwrites any existing configuration in /etc/yum.repos.d/
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
```

3. Install kubelet, kubeadm and kubectl:

```
yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

4. (Optional) Enable the kubelet service before running kubeadm:

```
sudo systemctl enable --now kubelet
```

```
Error: This command has to be run with superuser privileges (under the root user on most systems).
[ec2-user@ip-172-31-3-16 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes
Last metadata expiration check: 0:00:02 ago on Sat Sep 14 17:47:29 2024.
Dependencies resolved.
5.7 kB/s | 9.4 kB 00:01
```

| Package                  | Architecture | Version               | Repository  | Size  |
|--------------------------|--------------|-----------------------|-------------|-------|
| Installing:              |              |                       |             |       |
| kubeadm                  | x86_64       | 1.31.1-150500.1.1     | kubernetes  | 11 M  |
| kubectl                  | x86_64       | 1.31.1-150500.1.1     | kubernetes  | 11 M  |
| kubelet                  | x86_64       | 1.31.1-150500.1.1     | kubernetes  | 15 M  |
| Installing dependencies: |              |                       |             |       |
| conntrack-tools          | x86_64       | 1.4.6-2.amzn2023.0.2  | amazonlinux | 208 k |
| cri-tools                | x86_64       | 1.31.1-150500.1.1     | kubernetes  | 6.9 M |
| kubernetes-cni           | x86_64       | 1.5.1-150500.1.1      | kubernetes  | 7.1 M |
| libnetfilter_cthelper    | x86_64       | 1.0.0-21.amzn2023.0.2 | amazonlinux | 24 k  |
| libnetfilter_cttimeout   | x86_64       | 1.0.0-19.amzn2023.0.2 | amazonlinux | 24 k  |
| libnetfilter_queue       | x86_64       | 1.0.5-2.amzn2023.0.2  | amazonlinux | 30 k  |

```
Transaction Summary
Install 9 Packages

Total download size: 51 M
Installed size: 269 M
Downloading Packages:
(1/9): libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64.rpm 499 kB/s | 24 kB 00:00
(2/9): libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64.rpm 376 kB/s | 24 kB 00:00
(3/9): libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64.rpm 1.6 MB/s | 30 kB 00:00
(4/9): conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64.rpm 1.7 MB/s | 208 kB 00:00
(5/9): cri-tools-1.31.1-150500.1.1.x86_64.rpm 15 MB/s | 6.9 MB 00:00
(6/9): kubeadm-1.31.1-150500.1.1.x86_64.rpm 21 MB/s | 11 MB 00:00
(7/9): kubectl-1.31.1-150500.1.1.x86_64.rpm 17 MB/s | 11 MB 00:00
(8/9): kubernetes-cni-1.5.1-150500.1.1.x86_64.rpm 21 MB/s | 7.1 MB 00:00
(9/9): kubelet-1.31.1-150500.1.1.x86_64.rpm 29 MB/s | 15 MB 00:00

Total 45 MB/s | 51 MB 00:01
Kubernetes 8.0 kB/s | 1.7 kB 00:00
Importing GPG key 0x9A296436:
 Userid : "isv:kubernetes OBS Project <isv:kubernetes@build.opensuse.org>"
 Fingerprint: DE15 B144 86CD 3778 9E87 6E1A 2346 54DA 9A29 6436
 From : https://pkgs.k8s.io/core/stable/v1.31/rpm/repodata/repomd.xml.key
Key imported successfully
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
 Preparing : kubernetes-cni-1.5.1-150500.1.1.x86_64 1/1
 Installing : kubelet-1.31.1-150500.1.1.x86_64 1/9
 Installing : cri-tools-1.31.1-150500.1.1.x86_64 2/9
 Installing : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 3/9
 Running transaction
 Preparing : kubernetes-cni-1.5.1-150500.1.1.x86_64 1/1
 Installing : kubelet-1.31.1-150500.1.1.x86_64 1/9
 Installing : cri-tools-1.31.1-150500.1.1.x86_64 2/9
 Installing : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 3/9
 Installing : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 4/9
 Installing : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 5/9
 Installing : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 6/9
 Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 6/9
 Installing : kubelet-1.31.1-150500.1.1.x86_64 7/9
 Running scriptlet: kubelet-1.31.1-150500.1.1.x86_64 7/9
 Installing : kubeadm-1.31.1-150500.1.1.x86_64 8/9
 Installing : kubectl-1.31.1-150500.1.1.x86_64 9/9
 Running scriptlet: kubectl-1.31.1-150500.1.1.x86_64 9/9
 Verifying : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 1/9
 Verifying : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 2/9
 Verifying : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 3/9
 Verifying : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 4/9
 Verifying : cri-tools-1.31.1-150500.1.1.x86_64 5/9
 Verifying : kubeadm-1.31.1-150500.1.1.x86_64 6/9
 Verifying : kubectl-1.31.1-150500.1.1.x86_64 7/9
 Verifying : kubelet-1.31.1-150500.1.1.x86_64 8/9
 Verifying : kubernetes-cni-1.5.1-150500.1.1.x86_64 9/9

Installed:
 conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64 cri-tools-1.31.1-150500.1.1.x86_64 kubeadm-1.31.1-150500.1.1.x86_64
 kubectl-1.31.1-150500.1.1.x86_64 kubelet-1.31.1-150500.1.1.x86_64 kubernetes-cni-1.5.1-150500.1.1.x86_64
 libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[ec2-user@ip-172-31-3-16 docker]$
```

9. After installing Kubernetes, we need to configure internet options to allow bridging.

```
sudo swapoff -a echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf sudo sysctl -p
```

```
[ec2-user@ip-172-31-3-16 docker]$ sudo swapoff -a
[ec2-user@ip-172-31-3-16 docker]$ echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
net.bridge.bridge-nf-call-iptables=1
[ec2-user@ip-172-31-3-16 docker]$ sudo sysctl -p
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-3-16 docker]$
```

10. Initializing kubecuster:

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.3.16:6443 --token ekhyop.xkge2agz07jxxqqs \
--discovery-token-ca-cert-hash sha256:8206263b4e2632eb03dafa4819c7c8505d47b21e8ba8c4901d5802c791c806f7
[ec2-user@ip-172-31-3-16 docker]$ |
```

11. The mkdir command that is generated after initialization has to be copy pasted in the terminal.

```
[ec2-user@ip-172-31-3-16 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-3-16 docker]$
```

12. Then, add a common networking plugin called

flannel: kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[ec2-user@ip-172-31-3-16 docker]$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
[ec2-user@ip-172-31-3-16 docker]$ |
```

13. Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply this deployment file using this command to create a deployment kubectl apply -f

```
[ec2-user@ip-172-31-3-16 docker]$ kubectl apply -f https://k8s.io/examples/application/deployment.yaml
deployment.apps/nginx-deployment created
[ec2-user@ip-172-31-3-16 docker]$
```

<https://k8s.io/examples/application/deployment.yaml>

14. Use kubectl get pods to check if the pod is working correctly.

```
[ec2-user@ip-172-31-3-16 docker]$ kubectl get pods
NAME READY STATUS RESTARTS AGE
nginx-deployment-d556bf558-mvnj7 0/1 Pending 0 18s
nginx-deployment-d556bf558-w2pd8 0/1 Pending 0 18s
[ec2-user@ip-172-31-3-16 docker]$
```

15. To change status from pending to running use the following command:  
**kubectl describe pod nginx.**

```

nginx-deployment-d556bf558-w2pd8 0/1 Pending 0 18s
[ec2-user@ip-172-31-3-16 docker]$ kubectl describe pod nginx
Name: nginx-deployment-d556bf558-mvnj7
Namespace: default
Priority: 0
Service Account: <none>
Node: <none>
Labels: app=nginx
 pod-template-hash=d556bf558
Annotations: <none>
Status: Pending
IP: <none>
IPs: <none>
Controlled By: ReplicaSet/nginx-deployment-d556bf558
Containers:
 nginx:
 Image: nginx:1.14.2
 Port: 80/TCP
 Host Port: 0/TCP
 Environment: <none>
 Mounts: /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-8cms7 (ro)
Conditions:
 Type Status
 PodScheduled False
Volumes:
 kube-api-access-8cms7:
 Type: Projected (a volume that contains injected data from multiple sources)
 TokenExpirationSeconds: 3607
 ConfigMapName: kube-root-ca.crt
 ConfigMapOptional: <nil>
 DownwardAPI: true
 QoS Class: BestEffort
 Node-Selectors: <none>
 Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
 Type Reason Age From Message
 ---- -
Warning FailedScheduling 57s default-scheduler 0/1 nodes are available: 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }. preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.

Name: nginx-deployment-d556bf558-w2pd8
Namespace: default
Priority: 0
Service Account: default
Node: <none>
Labels: app=nginx
 pod-template-hash=d556bf558
Annotations: <none>
Status: Pending
IP: <none>
IPs: <none>
Controlled By: ReplicaSet/nginx-deployment-d556bf558
Containers:
 nginx:
 Image: nginx:1.14.2
 Port: 80/TCP
 Host Port: 0/TCP
 Environment: <none>
 Mounts: /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-6f18b (ro)
Conditions:
 Type Status
 PodScheduled False
Volumes:
 kube-api-access-6f18b:
 Type: Projected (a volume that contains injected data from multiple sources)
 TokenExpirationSeconds: 3607
 ConfigMapName: kube-root-ca.crt
 ConfigMapOptional: <nil>
 DownwardAPI: true
 QoS Class: BestEffort
 Node-Selectors: <none>
 Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
 Type Reason Age From Message
 ---- -
Warning FailedScheduling 57s default-scheduler 0/1 nodes are available: 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }. preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.

```

Use the below command to remove taints.

```

[ec2-user@ip-172-31-3-16 docker]$ kubectl taint nodes --all node-role.kubernetes.io/control-plane-
node/ip-172-31-3-16.ap-southeast-2.compute.internal untainted

```

16. Check the pod status.

```

NAME READY STATUS RESTARTS AGE
nginx 1/1 Running 1 (6s ago) 90s

```

17. port forward the deployment to your localhost so that you can view it

```
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

---

18. Verify your deployment Open up a new terminal and ssh to your EC2 instance. Then, use this curl command to check if the Nginx server is running.  
curl --head <http://127.0.0.1:8080>

Conclusion: In this experiment, we launched an EC2 instance and configured SSH access by updating the inbound rules. Next, we installed Docker and Kubernetes, and adjusted network settings to enable bridging. After completing the setup, we installed the Flannel networking plugin to ensure proper communication within the cluster. Once the cluster was up and running, we successfully deployed an NGINX server and verified its deployment.